

[Return to Module Overview Page \(https://onlinelearning.berkeley.edu/courses/665040/pages/module-seven\)](https://onlinelearning.berkeley.edu/courses/665040/pages/module-seven)

7.2: Exception Handling

Imagine this: you're writing a class that represents a database of students and their grades, schedules, etc. The constructor for this class accepts the name of a disk file as an argument; it reads the student information out of the file. What should the constructor do if the file does not exist? Remember, a constructor has no return value.

One possible solution is to set a special boolean "dead" member variable inside the object. Each of the object's member functions would then have to check the "dead" member before doing anything; if the flag was true, they could return a error value. For example, if the database class had an `n_students()` member, it could return -1 if the object was dead. Of course all the code that calls `n_students()` would need to check to see if it returned -1, and if so, it would have to handle things in some reasonable way. Whether the caller checked the error return value or not, though, at least the database class would always behave correctly. This solution is workable, but it adds a lot of essentially useless error-handling code to the application and to the database class.

Another possibility would be to set a dead flag, as above, in the database class, and trust the code that creates the database object to immediately discard the database if it turned out to be dead. This simplifies the code in the database class, because the individual methods don't need to check the dead flag, and in the application, because no dead database object should ever be used. But the word "trust" always should send your eyebrows shooting up. A class should never trust its caller to always do the right thing, as we've discussed in the past. A truly robust class doesn't trust code outside the class to do anything right!

The right answer in this case is neither of the two solutions above. Java provides a far more elegant solution: the use of exceptions.

```
import java.io.FileNotFoundException;

...

Database db = null;
try
{
    db = new Database("Students.dbf");
}
catch (FileNotFoundException fnfe)
{
    System.out.println("Database not found: "
                      + fnfe);
    System.out.println("Using default database");
    db = new Database("Default.dbf");
}
System.out.println("Database Initialized");

...
```

```
public class Database
{
    public Database(String filename)
        throws FileNotFoundException
    {
        if (!loadFile(filename))
            throw new FileNotFoundException("Cannot
                                           load
                                           Database");
    }
    // lots of code omitted, including
    //boolean loadFile(String name)
}
```

There are a lot of things going on here, so we'll look at it one step at a time, in the order that the code might execute.

The **try** keyword means precisely that: **try** to do something that might fail. In this case, we're going to **try** to create a **Database** object from the file `Students.dbf`. The flow of control transfers to the constructor function in class **Database**. Note that what follows the **try** keyword must be a block enclosed in curly brackets, even though it is just a single line of code. Note also that we must first declare the **db** variable outside of the **try** block if we are going to use it later in the program, so that it is not local to the **try** block!

In the **Database** constructor, we call a member function `loadFile()`, which presumably returns **true** on success. Let us first imagine that `loadFile` returns true. In this case, the `if()` test does not succeed, and the constructor returns normally. This puts us back in the mainline code, which assigns the newly constructed **Database** to the **db** variable. Flow of control proceeds then to the line that prints the "Database initialized" message.

But if instead `loadFile()` returns **false**, something very different happens. In this case, the `if()` succeeds, and we find ourselves at the line of code beginning with **throw**. We see a `java.io.FileNotFoundException` object being created (it's just an ordinary class object). What becomes of this object?

It helps to envision Robinson Crusoe stranded on his desert island. His only hope of communicating with the outside world is to write a message on a scrap of paper, stuff it into a bottle, and throw the bottle into the ocean. That is exactly what is happening here. By creating an exception object and throwing it, the constructor is sending a message to any code that advertises an interest in hearing about what went wrong. The Java runtime system takes that exception object and looks in the function that called this function for a **catch** block. A **catch** block is essentially a target for messages in bottles that might be thrown by code in trouble. The **catch** block in the mainline code above specifies that it is only interested in `FileNotFoundException` messages (actually, it would also receive subclasses of this exception type, if there were any). If there were not such a **catch** block in this function, Java would look in the function that called this function, and so on, until either such a block was found or until there were no more functions to check. (When that happens, the program is halted and a stack trace is printed out.)

A **catch** block must appear directly after the **try** block in which the code that might fail occurs. You can specify a number of **catch** blocks, one after another, all associated with a single **try** block. Each catch

block must specify a different exception type.

Notice that the `catch` block looks a little like a function definition:

```
catch (FileNotFoundException fnfe) { ... }  
  
int ketch (SomeOtherObjectType soot) { ... }
```

The resemblance is functional as well as visual. When the code in the `catch` block's pair of curly braces is executed, the variable `fnfe` has been initialized to refer to the thrown `FileNotFoundException` object. It's just like a function's formal argument list, except a `catch` block always has only one argument. Notice how the catch block above prints out the `FileNotFoundException` object.

A few more words about syntax are in order. First, a method must declare all the exceptions that it can throw, just as the `Database` constructor declares that it can throw `FileNotFoundException`. The compiler enforces this.

Exception classes such as `FileNotFoundException` are ordinary Java classes, although to be thrown, a class must inherit from `java.lang.Throwable`. Exceptions representing severe problems generally inherit from `java.lang.Error`, while those representing somewhat less extraordinary conditions inherit from `java.lang.Exception`. Both `Error` and `Exception` are subclasses of `Throwable`. If you define your own exception classes (you are encouraged to), they should generally be subclasses of `java.lang.Exception`.

Exceptions can be placed in two broad categories: unchecked (those that inherit from `Error` or from `RuntimeException`, a subclass of `Exception`) and checked (everything else.)

A method must either catch all checked exceptions thrown by other methods it calls, or declare that it itself throws those same exception types. The compiler enforces this also. It's actually fairly helpful in this regard, because it tells you exactly which exception types you must deal with. Not catching an exception is precisely the same as throwing one. This is why I previously told you to declare that your `main()` routine throws `IOException`; having declared that it threw `IOException`, `main()` didn't need to bother to catch it when any other functions threw it. Almost all of the methods in the `java.io` package can throw `IOExceptions`, as you can see in the API documentation.

You are not required (by the compiler, anyway) to declare or to catch unchecked exceptions. This is primarily for pragmatic reasons, since the JVM itself can throw many kinds of unchecked exceptions (`ArrayIndexOutOfBoundsException`, `ClassCastException`, `NullPointerException`, `OutOfMemoryError`) during the execution of ordinary code. Some unchecked exceptions should nonetheless always be caught and handled; `NumberFormatException` is one example.

There is one more interesting wrinkle: the *finally* clause. A set of `try` and `catch` blocks can be followed by a `finally` block:

```
SomeResource sr = new SomeResource();  
try  
{  
    sr.manipulate();
```

```
    }  
    catch (Failure f)  
    {  
        System.out.println("Failed!");  
        return;  
    }  
    finally  
    {  
        sr.dispose();  
    }  
    System.out.println("Done!");
```

The code in the **finally** block is executed after one of the **try** or **catch** blocks runs to completion (only one may), *even if that block returns from the method!* In the above example, even though the **catch** block returns, the **sr.dispose()** call would be made. Note that only if the **try** block runs to completion will the message "Done!" be printed. **finally** clauses let you put cleanup code that must be executed regardless of success or failure in one place, instead of duplicating it in each **try** and **catch** block.